

# Autonomous Driving through Imitation Learning

Ron Domingo

Department of Electrical Engineering  
Stanford University  
rdomingo@stanford.edu

**Abstract**—The goal of this paper is to train a policy for autonomous driving via imitation learning that is robust enough to keep the vehicle within the drivable area while avoiding obstacles that may be on the road surface. Using the Duckietown driving simulator, a vehicle is trained by an expert to drive around a collection of training tracks with both static and dynamic obstacles. The model is then exposed to a variety of unseen test tracks to evaluate their ability to drive a complete lap of each track without colliding into any obstacles or veering away from the roads. While a highly proficient driving model was not generated due to limited expert training data, the potential for improvement remains with future iterations of the model.

## I. INTRODUCTION

Path planning is a complex yet essential aspect of autonomous driving that is required for vehicles to traverse the world around us. Finding paths is a complicated task that involves the avoidance of both static and dynamic obstacles, all while ensuring that the vehicle stays within the drivable area of the road.

Imitation learning has become a popular method through which autonomous vehicles (AVs) are taught to handle the world around them. Given the complexity of the self-driving task, traditional reinforcement learning for the self-driving problem has proven to become intractable for the exponentially large state space associated with autonomous driving. Thus imitation learning can significantly speed up the learning process for AVs by embedding some proper behavior within the learning process of the vehicle [1].

This paper aims to apply imitation learning algorithms to teach an AV to drive around a closed loop track while avoiding obstacles. The AV will be simulated using the Duckietown [4] self-driving car simulator environments for OpenAI Gym. The goal of this paper is not only to drive the simulated vehicle around Duckietown without colliding into obstacles but also to ensure that the vehicle stays within the confines of the road surface, even when avoiding obstacles. The AVs will be evaluated on their ability to complete laps around various test tracks, where a full lap of the track without collision and off-road driving is considered a success.

For the purposes of the project, general road rules will not be considered when driving around the map and purely a drivable path around the road will be considered.

## II. RELATED WORKS

Applying imitation learning for autonomous driving is not a novel concept, in fact, this technique has been popularized by the research community as a way to more efficiently solve the

autonomous driving problem, especially in comparison with reinforcement learning.

The paper by Eike Rehder and Stiller [5] shows imitation learning being applied to path planning in an attempt to achieve human-like path planning from observation data. This concept is further expanded upon in the paper by Hussein [6] that created a model to perform 3D navigation through imitation learning. Both papers apply the concepts of the imitation learning algorithm to the specific path planning aspect of the autonomous driving problem, however, do not yet deal with the controls of the actual vehicle itself.

Several papers have gone even further and applied models learned through imitation learning into the real world. The papers by Pan [7] and Bansal [3] have shown that imitation learning has merit in creating an autonomous agent that can perform in the real world. Pan [7] demonstrates an autonomous vehicle that can drive around a high speed off-road track with a properly trained controller and Bansal [3] applies imitation learning to achieve well performing driving in real-world scenarios.

While this paper will not explicitly utilize the frameworks used in each of these papers, the literature provides ample examples on how imitation learning can be applied to teach autonomous vehicles how to properly drive in a more time efficient and less data intensive way. The concepts in these papers will help to inform how the model will be formed to create a self driving agent in the Duckietown simulator.

## III. SETUP

### A. Simulation Environment

The Duckietown self-driving car simulator is an OpenAI Gym extension that serves as a simulator for the Duckietown hardware robots. The driving environment features marked roads, intersections, static objects such as buildings and trees, and finally dynamic obstacles like pedestrians that move around throughout the maps randomly.



Fig. 1. Simulation environment screenshot.

The Duckietown simulator outputs observations from a monocular camera mounted on the robot with image sizes of  $640 \times 480 \times 3$  as seen in Figure 2.



Fig. 2. Vehicle observation example.

Duckietown features customizable tracks ranging from mini-cities complete with intersections, signage, pedestrians, and buildings, to simplified straight line tracks. For the purposes of this paper, custom closed loop tracks are considered since the AV will be evaluated on its ability to maneuver around the track and get back to its original position without colliding with any obstacles along the way. The tracks, however, will be customized so as to provide the AV with various training and test tracks for evaluation. Each track will also include a collection of randomly placed obstacles throughout, whether they are static or dynamic obstacles on or off the road surface. This is done to ensure that the AV can handle maneuvers such as stopping for pedestrians or swerving around static objects that lie in the road. Another purpose of this randomization is to prevent the AV from simply memorizing the tracks and different maneuvers from the testing tracks.

### B. Input Processing

The Duckietown simulator provides observations in terms of images of size  $640 \times 480 \times 3$ . The first step of image pre-processing is to downscale the images in order to improve

computational efficiency. The observations are downsampled to a fourth of their original size through nearest-neighbor resampling, generating images that are  $160 \times 120 \times 3$ . The downsampled images are then normalized before being passed to the Convolutional Neural Network (CNN).

### C. Action Generation

The simulation vehicle takes in two parameters for each action, the velocity and the steering angle, each of which take on values between -1 and 1. A CNN is used to convert the raw images to the desired action that would enable the vehicle to get around the track. The CNN architecture is as follows:

- Convolutional Block
- Convolutional Block
- Convolutional Block
- Convolutional Block
- Dropout
- Fully Connected Layer
- Fully Connected Layer




Fig. 3. Neural Network Diagram

In this architecture, each convolutional block consists of a convolution, rectified linear unit, and a batch normalization. The 4 convolutional blocks extract features from the AV camera while reducing the state space of the image. Then information is then passed to the 2 fully connected layers to convert the states to the action space.

While the actions accepted by the simulator are continuous, the network is trained with only 4 actions - left, right, forward, and stop. This is done to reduce the action space of the model and thereby reduce the computation required to perform the task at hand.

Two CNN models are tested throughout this paper. While both models share the same basic architecture as the one outlined above, the models differ in their input layer. The first model uses a single camera frame to generate an output for the vehicle, essentially utilizing an input of size  $160 \times 120 \times 3$ . For the purposes of this paper, this model will be referred to as the single-frame model. The second model, on the other hand, uses the current and the previous 3 frames to generate an action. These 4 frames are stacked together in order, resulting in an input size of  $160 \times 120 \times 12$ . The larger input is an attempt to preserve some temporal relationship between the frames. The

goal of the second model is to try and improve the driving performance given some temporal context for where the car is heading on the track over time. In the further sections of the paper, the second model will be referred to as the multi-frame model.

#### IV. IMITATION LEARNING

The learning process used in this paper follows the process outlined in the deep active imitation learning paradigm by Ross [8]. The process is as follows: (1) Collecting demonstrations. (2) Supervised training of the neural network. (3) Active learning to refine the initially learned policy.

##### A. Collecting Demonstrations

Imitation learning requires expert behavior examples in order to train the model. For this paper, the expert examples were obtained from manually driving multiple laps of each test track. During the laps, the observation state and the action taken were recorded.

Each lap has a randomized start position and orientation and a lap is complete when the vehicle has reached close to its starting position on the road. Only valid laps were recorded and stored for training purposes.

##### B. Supervised Training

Once recorded, the network is then trained on batched samples of the recorded observations and actions. The loss function used during the training process is an L2 loss function. With the single-frame model, batches of observations and the corresponding actions were simply taken for each training epoch. With the multi-frame model, preprocessing had to be done to the raw observation data. The preprocessing involved pre-stacking frames together in a sliding window fashion to generate a list of inputs, where each input were 4 frames stacked together, and their corresponding action was the action taken at the most current frame in the stack. Both models were trained with the same expert data and each model was trained for 100,000 epochs.

##### C. DAGGER Re-Training

After the supervised training has completed, each model is then run through training tracks again, but this time autonomously. While autonomously driving, the expert user will be logging the actions they would have taken for any given observation. These expert actions, as well as the observations for any given timestep, are logged without affecting the running model so that the expert actions are attached to the trajectories introduced by the model when it veers off course.

The models are then retrained with the new trajectories and actions to better improve the driving ability of the models in previously unseen trajectories.

## V. RESULTS

### A. Supervised Training

After training both models, the training loss graphs indicate some convergence over the training datasets. On first inspection, the multi-frame model, in Figure 5, converges much faster and with a smaller error compared to the single-frame model in Figure 4.




Fig. 4. Single-frame model training loss




Fig. 5. Multi-frame model training loss

### B. DAGGER Re-training

DAGGER re-training was performed over 50,000 epochs for either model. In the case of the single-frame model, in Figure 6, the model seemed to converge over the new DAGGER dataset, albeit with greater error than the training dataset.




Fig. 6. Single-frame model DAGGER re-training loss

This was not the case, however, for the multi-frame model in Figure 7. After DAGGER re-training, the model does not appear to converge over the new data and generates much larger errors than seen before, despite its low errors in the initial supervised training.




Fig. 7. Multi-frame model DAGGER re-training loss

### C. Testing

After the training process, the models were evaluated on their ability to complete laps of unknown test tracks without any collisions, all while staying on the drivable area of the road. The models were run autonomously on a variety of test tracks, each with their own distinct set of static and dynamic obstacles. For every failed or completed lap, the AV would respawn in a new random location with a random orientation. The success rate of the AVs is then recorded and used to evaluate the driving ability of the models.

1) *Single-Frame Model*: The single-frame model was able to achieve a 4.5% success rate when run over thousands of iterations on unknown tracks. After the DAGGER re-training, the performance of the model actually decreased and dropped to a 0% success rate.

2) *Multi-Frame Model*: The multi-frame model was not able to complete any successful laps when run on unknown tracks. Similarly, after the DAGGER re-training, the multi-frame model did not improve with the corrected data.

## VI. DISCUSSION

Overall, the two models did not perform well with unknown tracks. Most surprising is that the multi-frame model performed far worse than the single-frame model, despite a more promising training loss curve.

The difference between the two models can be seen when watching each model drive autonomously. The single-frame model is able to make decisions, some correct and some not, despite having never seen those particular observations. This model appeared to be able to turn and keep itself on the track rather well, but would fail in avoiding both static and dynamic obstacles on the road.

The multi-frame model, on the other hand, was prone to getting stuck when faced with unseen trajectories. In these situations, the multi-frame model would either stop completely or drive around in circles, trying to find a road to drive on.

The DAGGER re-training also did not help with these behaviors. The expert action collection during the DAGGER algorithm has to be improved so as to provide a more accurate correction to the learned policies. The data for the re-training was collected with a user on a keyboard watching the AV drive autonomously. Without actual control of the vehicle, the user was left to hitting keys that they felt like would correct the vehicle trajectory, but without actual feedback, the actions logged by the user could be mistimed and therefore incorrect.

The lack of convergence in the training loss given the new data indicates that the new actions varied significantly from the actual expert driving when the expert has full control of the vehicle. This is also shown by the consistent failures of the AVs when run on the post-DAGGER models.

For both models, a larger expert dataset to train on would prove extremely beneficial. The manual collection of data proved to be a limiting factor in the performance of the trained models as a more comprehensive dataset with more examples of obstacle avoidance could have helped with the single-frame model specifically. While the performance on unknown tracks was low with the single-frame model, its performance when evaluated on the same tracks it was trained on was much better at a 37% success rate. This, however, could be due to memorizing the routes on those particular tracks over the unknown tracks, although it does show some promise. The same performance can not be said for the multi-frame model on the training tracks as it still failed to complete any successful laps on previously seen tracks.

Perhaps the poor performance of the multi-frame model was due to the fact that the attempt to incorporate temporal information into the decision making was too naive. Instead, a full restructuring may be needed to properly incorporate the temporal aspects of motion planning into the AV. This could be done by implementing an LSTM network to better take advantage of temporal relationships between objects and the vehicle while driving around the track, as done by Bai [2].

## VII. CONCLUSION

Path planning is an extremely challenging issue that must be solved before vehicles can become truly autonomous. There have been many attempts to try and solve this challenge through traditional or deep reinforcement learning techniques, but the complexity of the problem leads to time-consuming and tedious methods that become cumbersome to compute. Imitation learning has gained traction over the years as an alternative to reinforcement learning by allowing researchers to transfer human knowledge to an agent, thereby expediting the process of obtaining well performing policies without having to explicitly define an extremely complex model to approximate autonomous driving.

This paper applied the imitation learning algorithms to the Duckietown driving simulator to create an autonomous driving agent that could navigate a variety of tracks without any collisions all while staying within the road boundaries. Two models were tested: a model that accepted a single camera frame and outputted an action, and a model that took in the last 4 frames of driving to generate the best action for the vehicle. After training, the single-frame model fared much better than the multi-frame model, however, both models share room for significant improvement in future iterations of this project. The limited expert data greatly affected the driving performance of the models and the naive attempt to include temporal relationships in the action generation did not do much to improve performance. More work is required to create models that can accurately drive around the simulator with

competence, but the current results are still significant given that the models were trained without an explicit definition of the problem, its rewards, the state transitions, or any ground truth actions. In this way, this paper was still able to show the benefits of imitation learning in trying to solve high-dimensional decision-making problems.

#### REFERENCES

- [1] Sharone Dayan Alexandre Attia. Global overview of imitation learning. Technical report, 2018.
- [2] Zhengwei Bai. Deep learning based motion planning for autonomous vehicle using spatiotemporal lstm network. Technical report, 2019.
- [3] Mayank Bansal. Chauffeurnet: Learning to drive by imitating the best and synthesizing the worst. Technical report, 2018.
- [4] Maxime Chevalier-Boisvert, Florian Golemo, Yanjun Cao, Bhairav Mehta, and Liam Paull. Duckietown environments for openai gym. <https://github.com/duckietown/gym-duckietown>, 2018.
- [5] Jannik Quehl Eike Rehder and Christoph Stiller. Driving like a human: Imitation learning for path planning using convolutional neural networks. Technical report, 2017.
- [6] Ahmed Hussein. Deep imitation learning for 3d navigation tasks. *Neural Computing & Applications*, 29(7):389–404, 2018.
- [7] Yunpeng Pan. Agile autonomous driving using end-to-end deep imitation learning. Technical report, 2019.
- [8] Stéphane Ross. A reduction of imitation learning and structured prediction to no-regret online learning. Technical report, 2011.